



# Extending Applications Safely and Efficiently




Yusheng Zheng<sup>1</sup> • Tong Yu<sup>2</sup> • Yiwei Yang<sup>1</sup> • Yanpeng Hu<sup>3</sup>

Xiaozheng Lai<sup>4</sup> • Dan Williams<sup>5</sup> • Andi Quinn<sup>1</sup>

<sup>1</sup>UC Santa Cruz    <sup>2</sup>eunomia-bpf Community    <sup>3</sup>ShanghaiTech University

<sup>4</sup>South China University of Technology    <sup>5</sup>Virginia Tech

# Extensions are everywhere



## What are extensions?

- Customize software without modifying source code

## Why do we need them?

- Easier to maintain and update

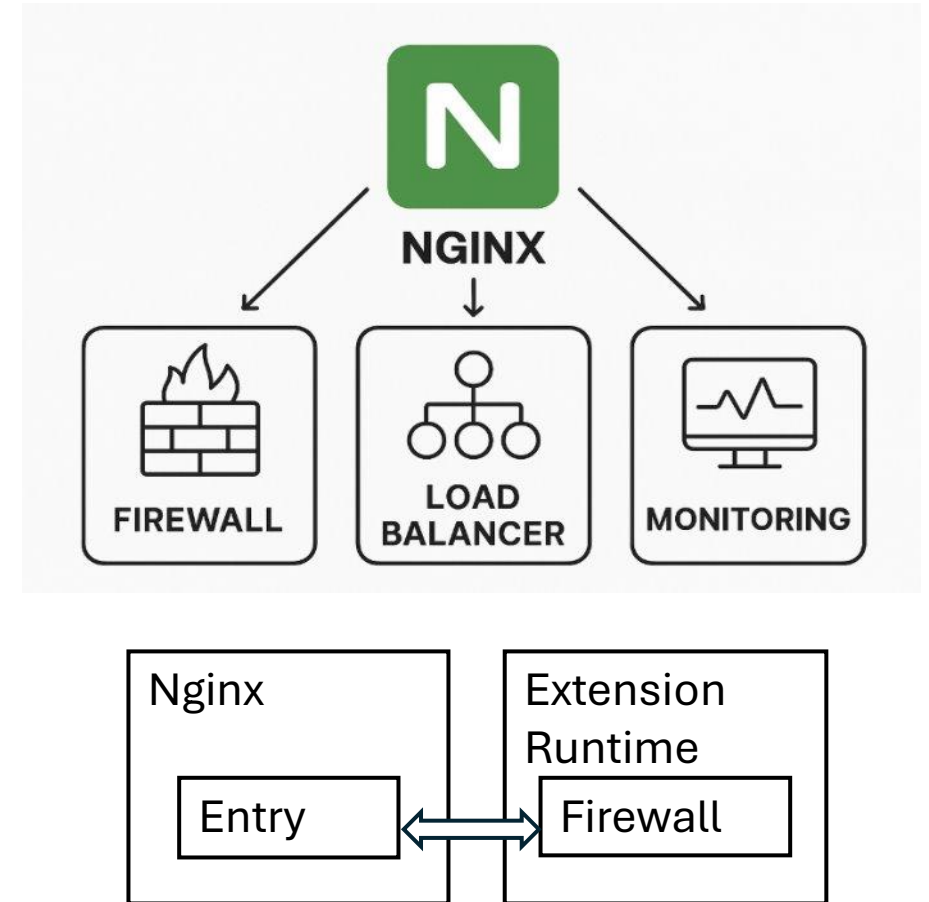
# Nginx firewall example

## Before deployment, user:

- Writes firewall using nginx APIs
- Associates firewall with request processing extension entry.

## During runtime, Nginx:

- Jumps to firewall when reaching request processing entry.
- Executes firewall in the extension runtime execution context.



# Extension Problems & requirements

## Real-world safety violations:

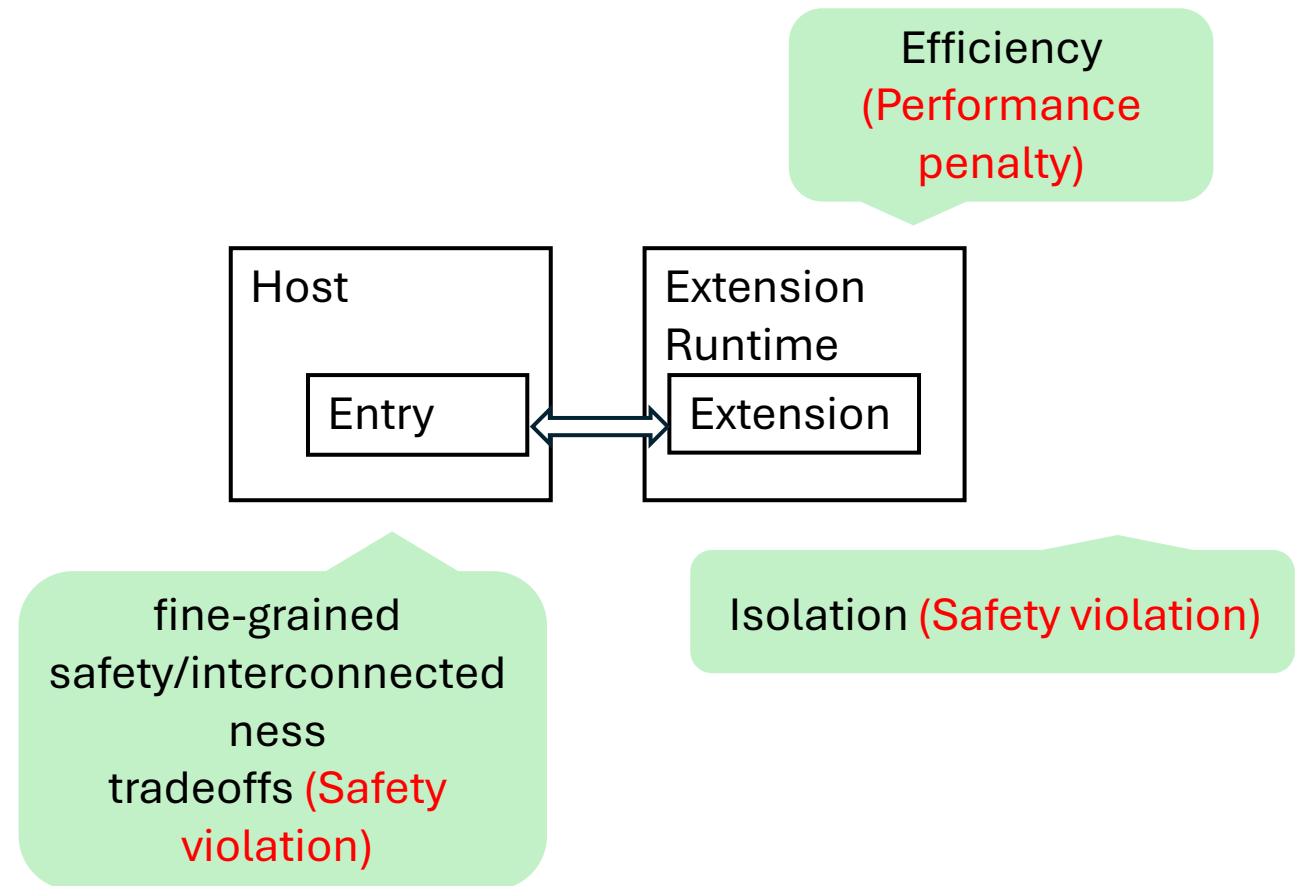
- Bilibili CDN outage, Apache buffer overflow, Redis RCE

## Performance penalty:

- WebAssembly/Lua impose 10-15% overhead

## Requirements:

- Fine-grained safety and interconnectedness trade-offs
- Isolation
- Efficiency



# State-of-the-Art Falls Short

- Dynamic loading: efficiency but no isolation or fine-grained safety-interconnectedness policies (LD\_PRELOAD, DBI tools)
- Software Fault Isolation: safety with 10–15 % performance penalty (XFI [OSDI 06], NaCL [SOSP 09], RL-Box [USENIX Security 20], Wasm and Lua)
- Subprocess: strong isolation but high IPC overhead (Wedge [NSDI 08], Shreds [IEEE SP 16], lwC [OSDI 16], and Orbit [OSDI 22])
- Kernel eBPF uprobes: isolation at micro second-level trap cost, low efficiency

# Contributions

## **Extension Interface model (EIM)**

Navigate fine-grained  
safety/interconnectedness trade-offs for  
extensions

## **Bpftime**

Efficient support for EIM and isolation  
through userspace eBPF runtime

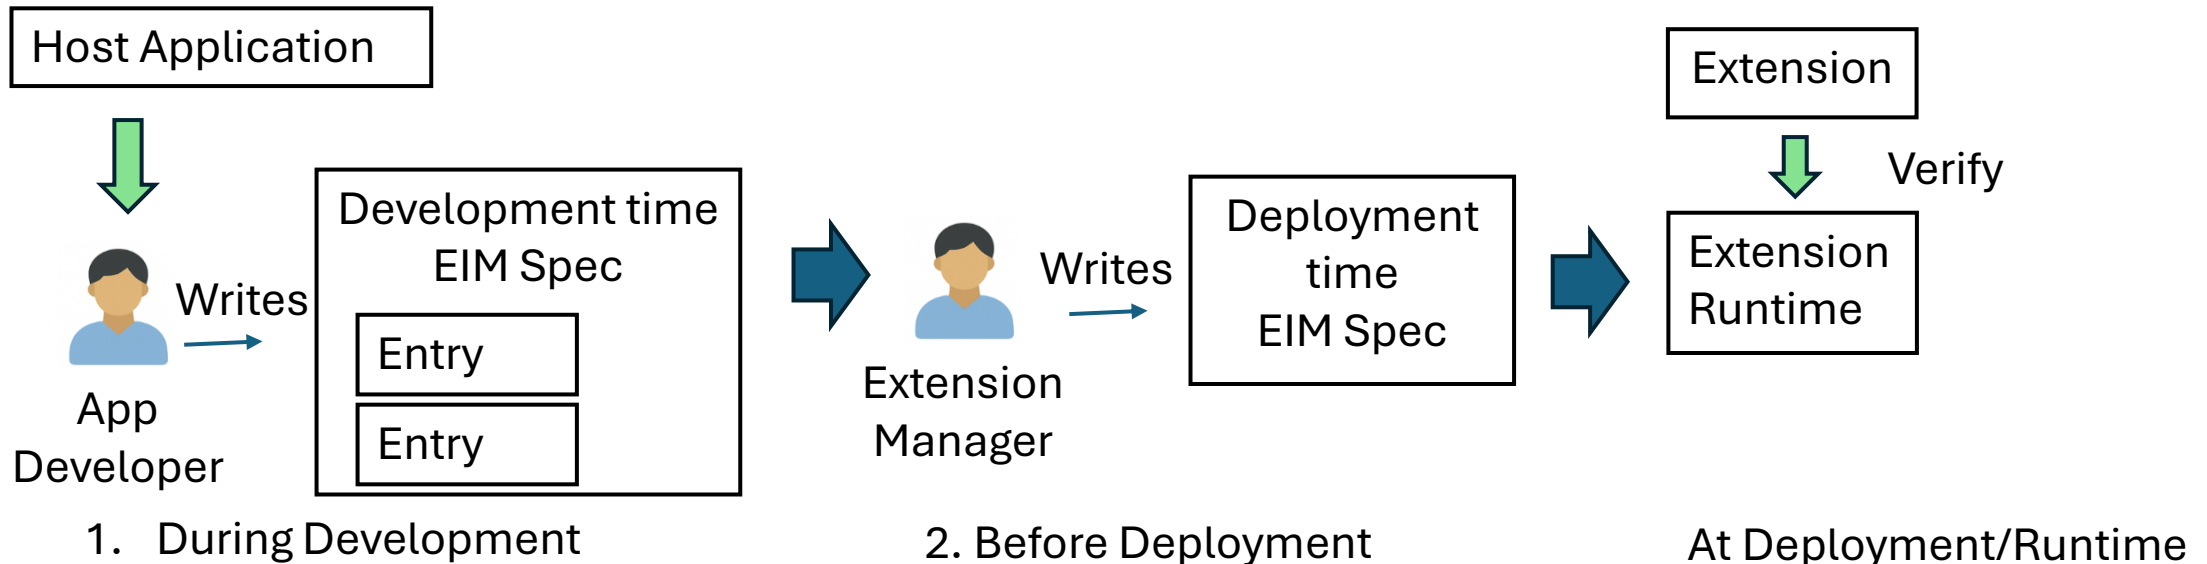
Up to 6x less overhead than current state-of-the-art!

# Outline

- Background & motivation: Extensions
- → **Extension Interface Model (EIM)**: Fine-grained Interface
- bpftime Runtime: safety & performance
- Evaluation

# EIM: Extension Interface Model

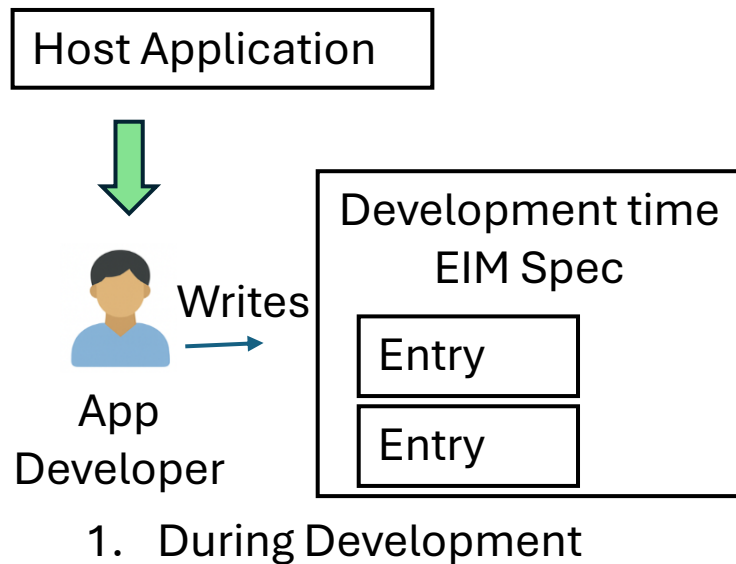
- Challenge:
  - Enable fine-grained safety/interconnectedness trade-offs
- Solution:
  - Two-phase specification (Development Time and Deployment Time)
  - Model all resources as capabilities





# EIM: Development Time Specification

- Developers annotate code for capabilities
- Automatically extracted into capability manifest



---

```
EIM_STATE_DEFINE(readPid, read, ngx_pid);  
EIM_HFUNC_DEFINE_WITH_CONSTRAINTS(  
    ngxTime,  
    HF_RET_POSITIVE);  
EIM_EXTENSION_ENTRY_DEFINE(  
    processBegin,  
    ngx_http_process_request,  
    int,  
    struct Request *);
```

---

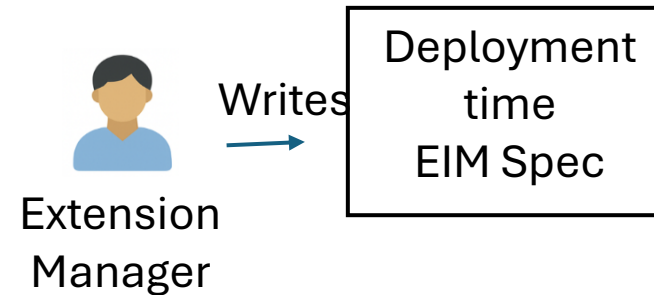
# EIM: Deployment Time Specification

- YAML policies specify safety/interconnectedness tradeoffs
- Compact policies (avg of 30 lines in evaluation).

---

```
1 Extension_Class(  
2   name = "observeProcessBegin",  
3   extension_entry = "processBegin",  
4   allowed = {instructions<inf, nginxTime,  
               readPid, read(r)})  
5 Extension_Class(  
6   name = "updateResponse",  
7   extension_entry = "updateResponseContent"  
8   allowed = {instructions<inf, read(r),  
               write(r)})
```

---



## 2. Before Deployment

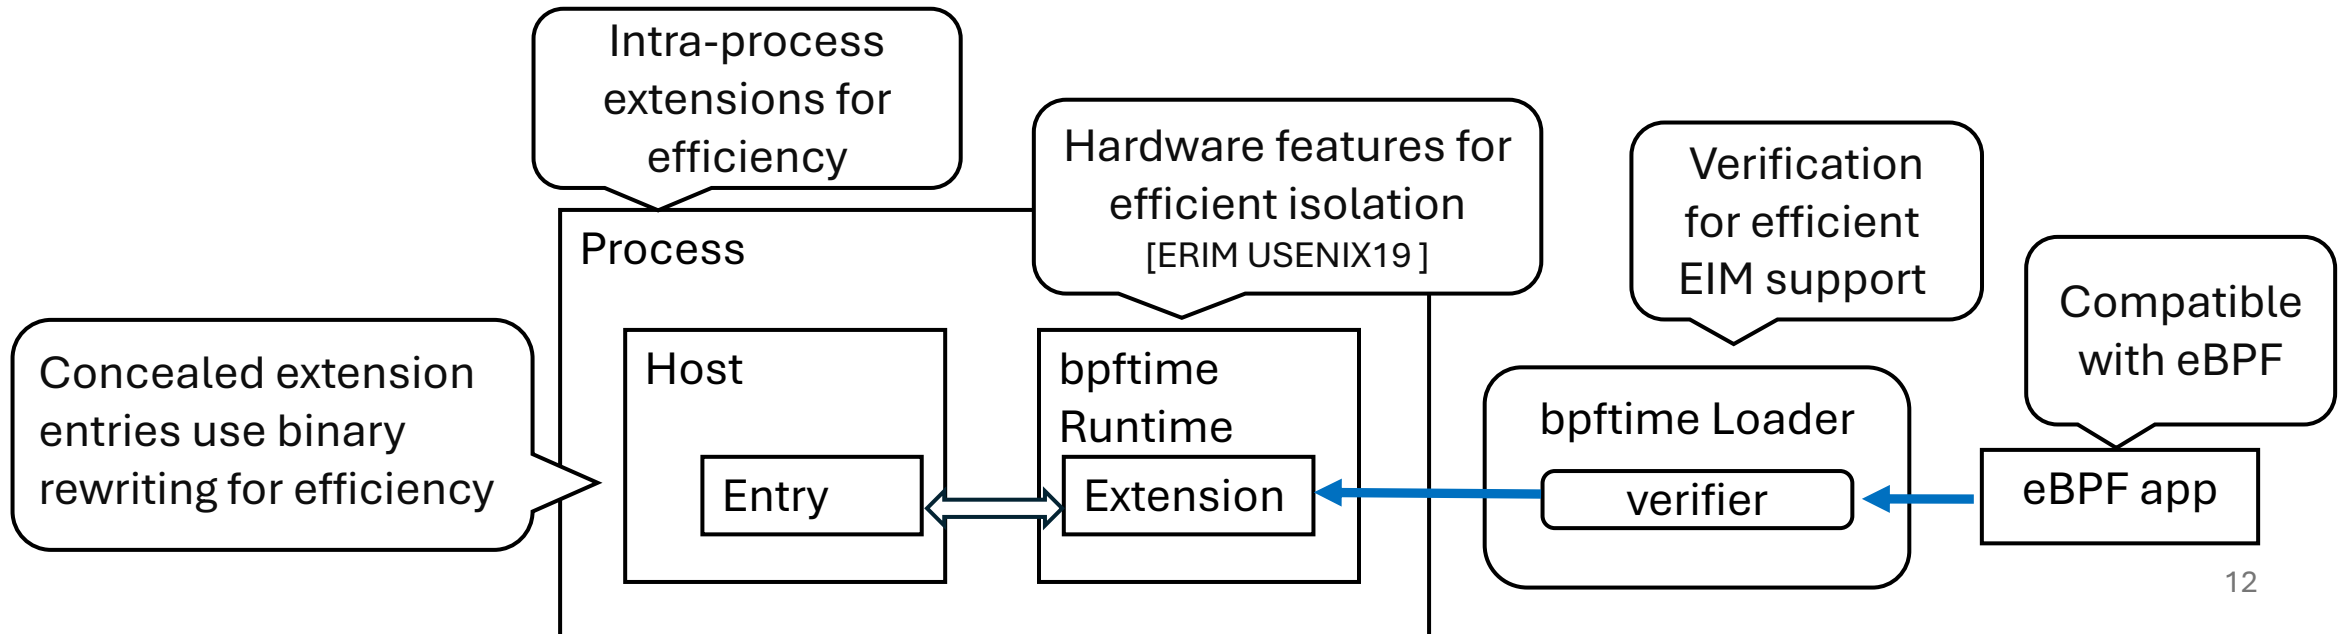
-

# Outline

- Background & motivation: Extensions
- Extension Interface Model (EIM): Fine-grained Interface
- → **bpftime Runtime**: safety & performance
- Evaluation

# bpftime: userspace eBPF extension framework

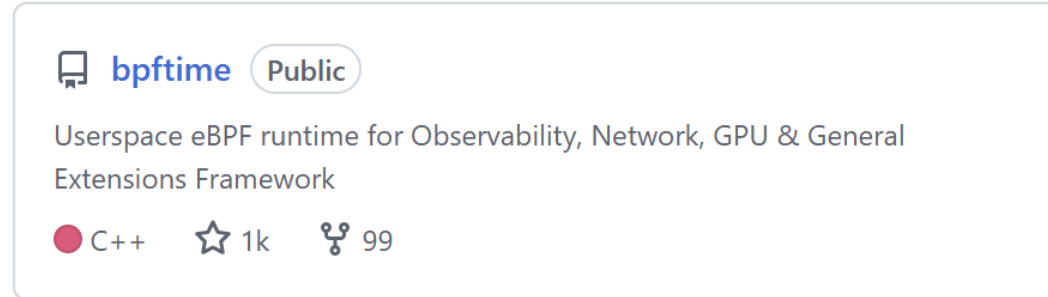
- Challenge:
  - efficiently support EIM and isolation
- Solution:
  - exploit eBPF-style verification, binary rewriting, and hardware features



# Outline

- Background & motivation: Extensions
- Extension Interface Model (EIM): Fine-grained Interface
- bpftime Runtime: safety & performance
- → **Evaluation**

# Six Real-World Use Cases



GitHub: <https://github.com/eunomia-bpf/bpftime>

## Customization

- Nginx Firewall
- Redis Durability
- FUSE Metadata Cache


## Observability

- DeepFlow
- Syscount
- Sslsniff

# Customization: Nginx firewall

- 5× to 6× less overhead than lua or WebAssembly

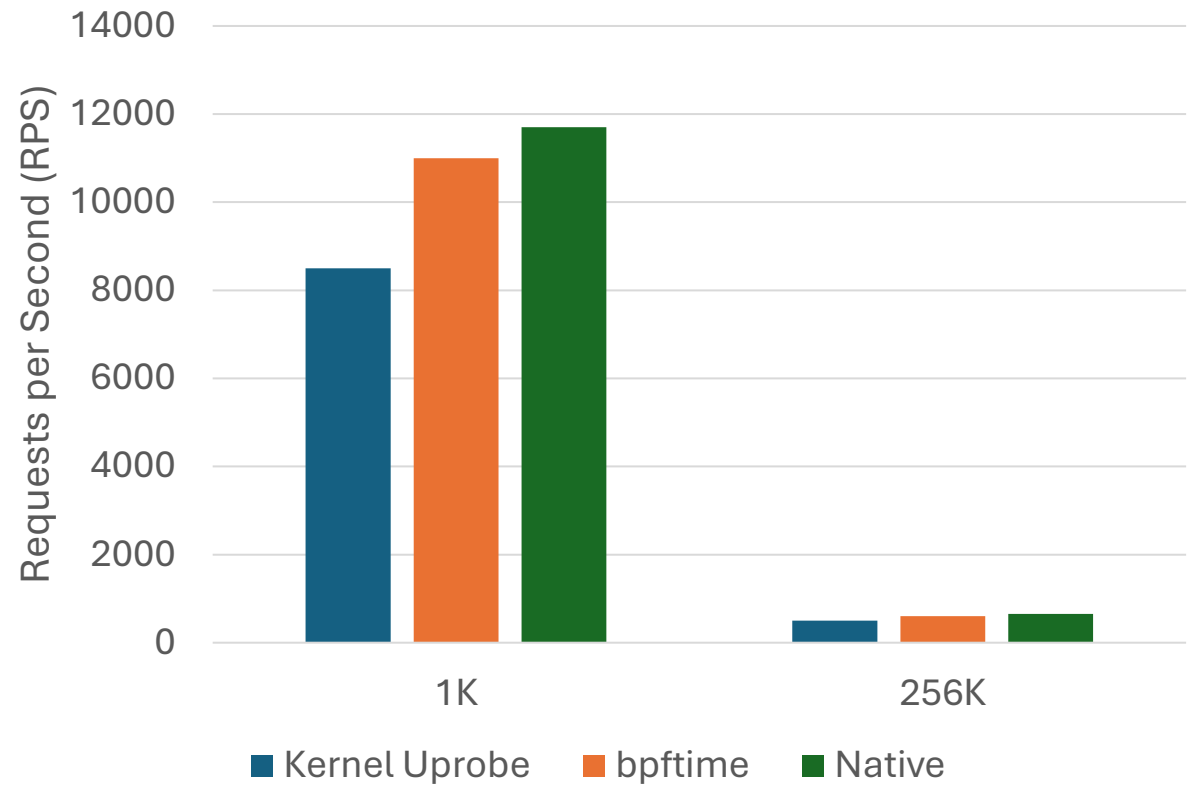
better



# Observability: sslsniff

- Maximum 21% less overhead than kernel eBPF

better





# Contributions

# Questions?

## Extension Interface model (EIM)

Navigate fine-grained  
safety/interconnectedness trade-offs for  
extensions

## Bpftime

Efficient support for EIM and isolation through  
userspace eBPF runtime

Up to 6x less overhead than current state-of-the-art!



```
bpftime load ./example/malloc/malloc
bpftime start nginx -c ./nginx.conf
```



# Backup

# bpftime: userspace eBPF extension framework

- Challenge for compatibility and efficiency:
  - eBPF: tightly coupled components
  - Bpftime: Intercept syscalls & Share memory maps

